

Performance of rdt3.0



- rdt3.0 correct, but performance stinks
- Calculate:
 - Transmission delay of 8000bit packet on 1Gbps link

$$D_t x = rac{L}{R} = rac{8000b}{10^9 bps} = 0.008 ms$$

Link utilization, or the fraction of time sender is transmitting,
assuming 15ms propagation delay and negligible size ACK packets

$$U_{snd} = rac{L/R}{RTT + L/R} = rac{0.008}{30.008} = 0.00027$$

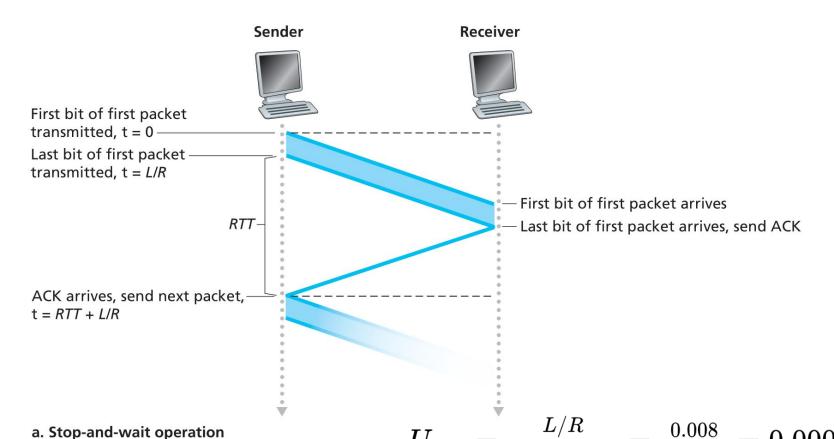
- Link goodput, or rate of data delivery over time

$$8000b/30.008ms = 267kbps$$

On a 1Gbps link!

rdt3.0: stop-and-wait operation



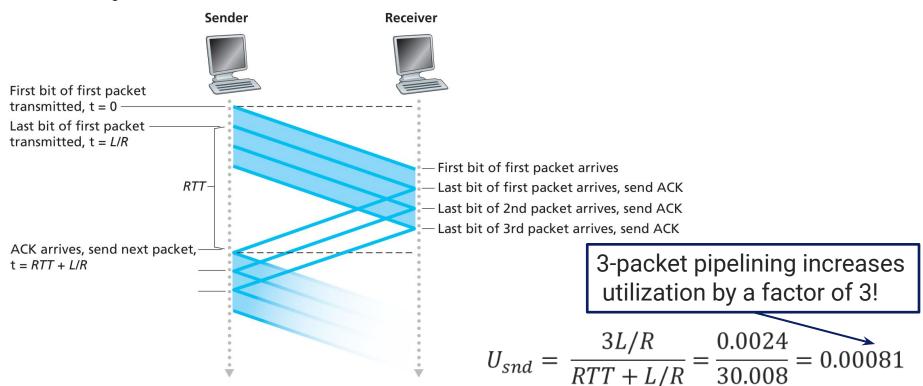


Pipelining: increased utilization



Pipelining: Sender allows multiple, "in-flight", yet-to-be-acknowledged packets

- Range of sequence numbers must be increased
- Buffering at sender and/or receiver



Pipelined protocols: Overview



Go-back-N (GBN)

- Sender can have up to N unacked packets in pipeline
- Receiver only sends cumulative ack
 - ACK for last contiguous packet
 - No ACK for new packets past a sequence number gap
- Sender has timer for oldest unacked packet
 - When timer expires, retransmit all unacked packets

Selective Repeat (SR)

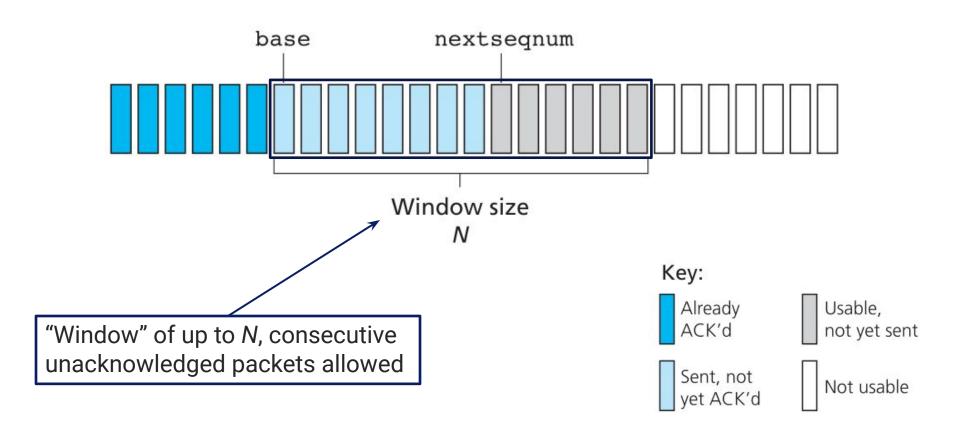
- Sender can have up to N unacked packets in pipeline
- Receiver sends individual ACK for each packet

- Sender maintains timer for each unacked packet
 - When timer expires, retransmit only that unacked packet

What are the relative advantages of each protocol?

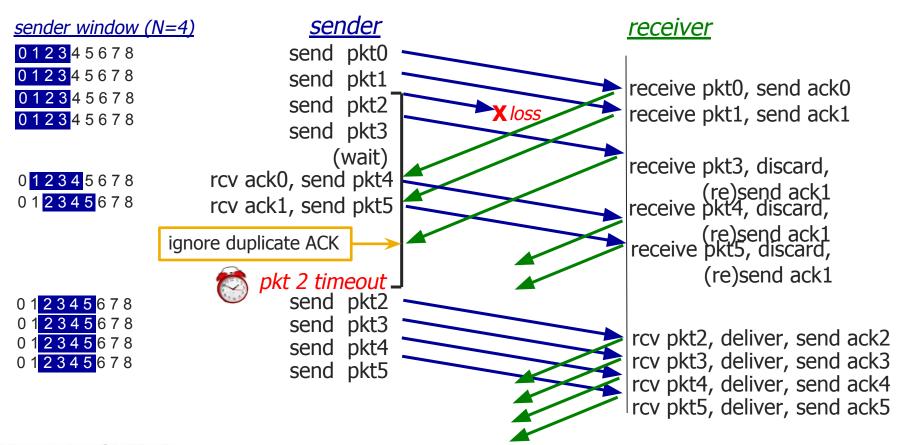
GBN: Sender window





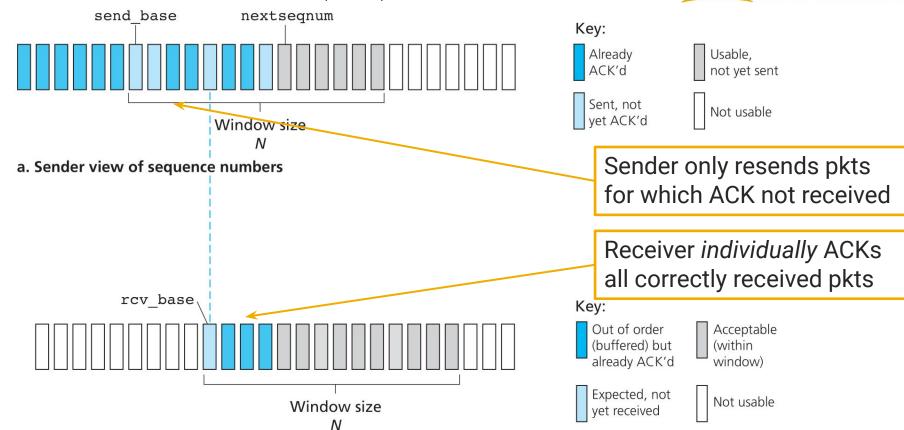
GBN in action





Selective Repeat (SR)

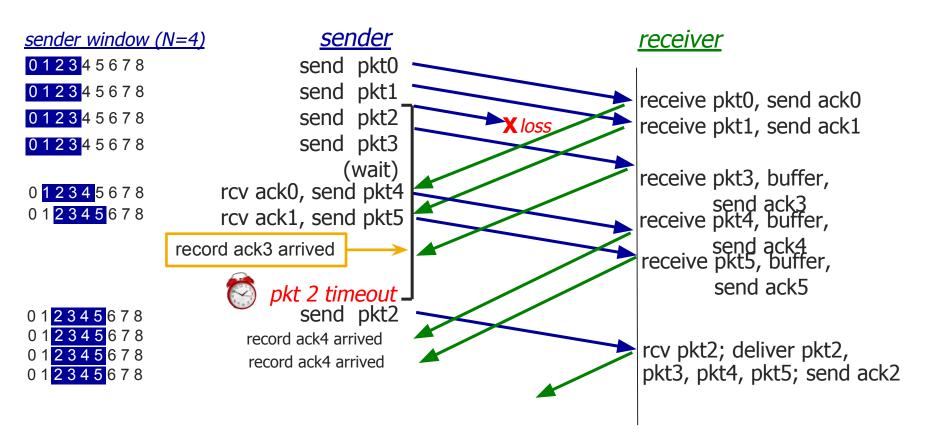




b. Receiver view of sequence numbers

SR in action

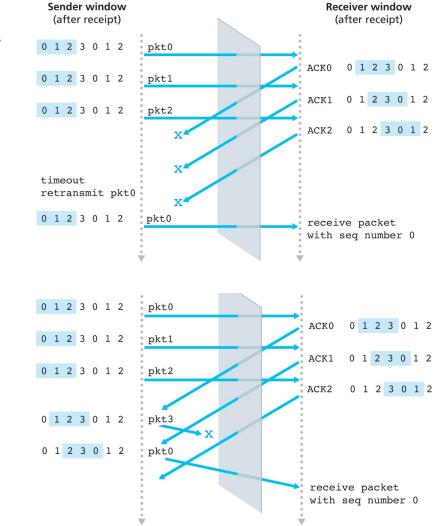




Selective Repeat Dilemma

- Receiver sees no difference in the two scenarios!
- Duplicate data accepted as new in the second scenario!

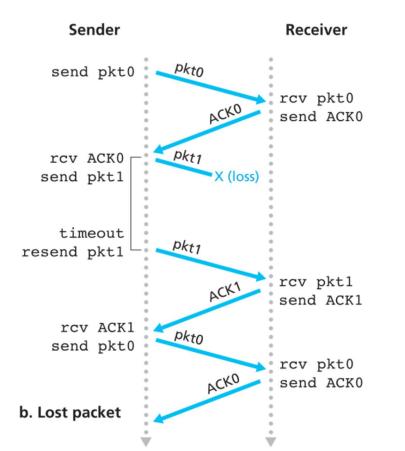
Q: What should be the relationship between window size and sequence numbers to avoid the problem?





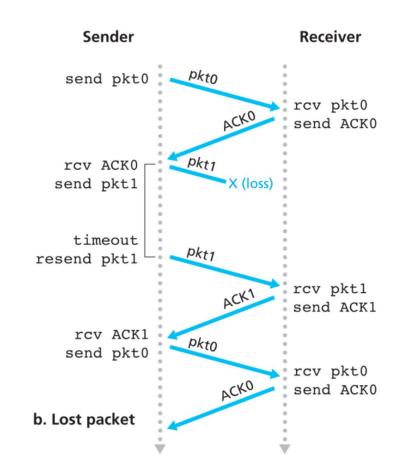
- Q: How to set TCP timeout value?
 - Too short: premature timeout, unnecessary retransmissions
 - Too long: slow reaction to segment loss

- Longer than RTT?
 - but RTT varies





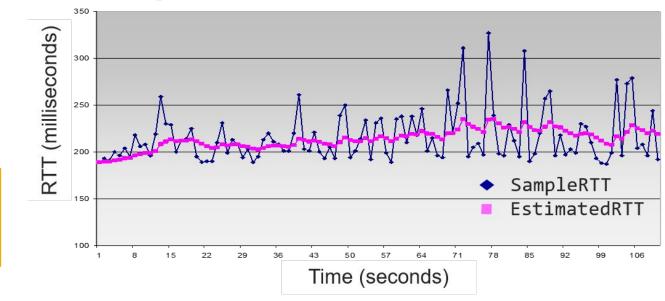
- Q: how to estimate RTT?
 - SampleRTT: measured time from segment transmission until ACK receipt
 - SampleRTT will vary, want estimated RTT "smoother"
 - Average several recent measurements, not just current SampleRTT





- Exponential weighted moving average
 - EstimatedRTT = $(1-\alpha)$ *EstimatedRTT + α *SampleRTT
 - Typical value: α =0.125

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr





- Timeout interval:
 - EstimatedRTT plus "safety margin"
 - large variation in EstimatedRTT → larger safety margin
- Estimate SampleRTT deviation from EstimatedRTT:

DevRTT =
$$(1-\beta)*DevRTT + \beta*|SampleRTT-EstimatedRTT|$$

TimeoutInterval = EstimatedRTT + 4*DevRTT

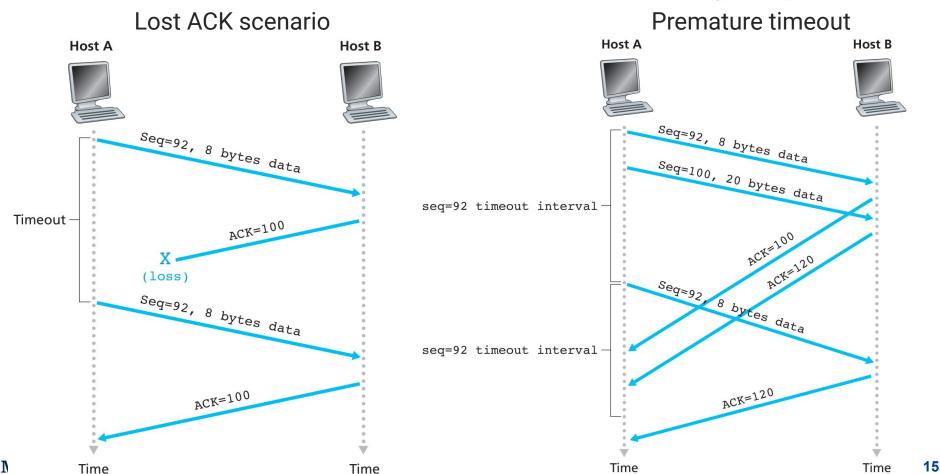


Estimated RTT

"safety margin"

TCP retransmission scenarios

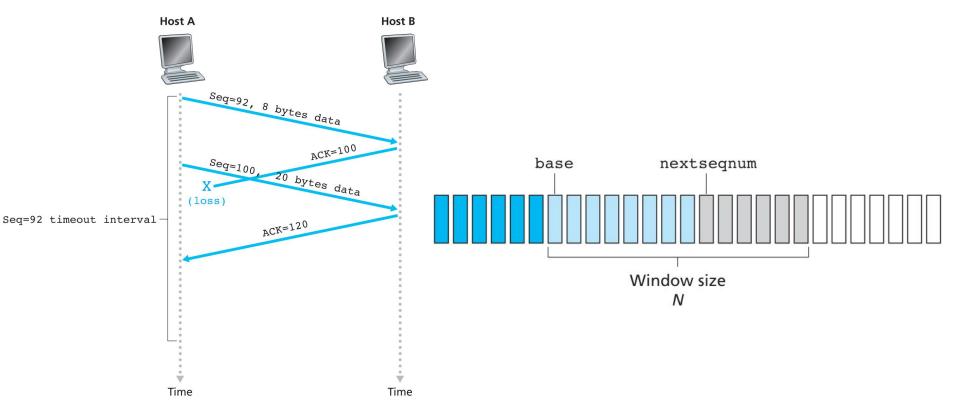




TCP: retransmission scenarios



Cumulative ACK



TCP ACK generation



Event at receiver	TCP receiver action
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send duplicate ACK, indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

TCP fast retransmit

MONTANA STATE UNIVERSITY

- Time-out period often relatively long before packet loss detected
- Detect lost segments via duplicate ACKs
 - Many packets in pipeline
 - If segment lost, many duplicate ACKs generated
- TCP fast retransmit
 - Resend packet after 3 duplicate ACKs
 - Do not wait for timeout packet probably already lost

