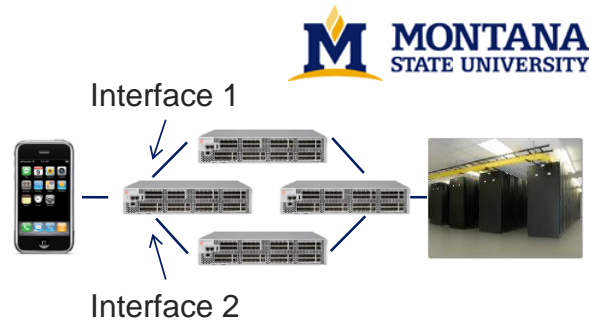


Datagram networks

- Datagram network provides network-layer *connectionless* service
 - No call setup at network layer
 - Routers: no state about end-to-end connections
- Packets forwarded using destination host address
 - How to keep routing tables small?
 - Longest prefix matching



Address range	Interface
128.11.52.0 - 128.11.52.255	1
153.90.2.0 - 153.90.2.255	2
153.90.2.87 - 153.90.2.89	1

Longest Prefix Matching



Routing Table

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Match + Action

Prefix Match	Link Interface
11001000 00010111 00010	0
11001000 00010111 00011000	1
11001000 00010111 00011	2
otherwise	3

Longest Prefix Matching



Routing Table

Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

Match + Action

Prefix Match	Link Interface
11001000 00010111 00010	0
11001000 00010111 00011000	1
11001000 00010111 00011	2
otherwise	3

Translate the routing table from binary to CIDR representation

200.23.16.0 to 200.23.23.255	→ 0
200.23.24.0 to 200.23.24.255	→ 1
200.23.25.0 to 200.23.31.255	→ 2
else	→ 3
200.23.16.0/21	→ 0
200.23.24.0/24	→ 1
200.23.24.0/21	→ 2
else	→ 3

201

Mountains & Minds



Chapter 5: Network Layer – Control Plane


Mountains & Minds

202

Introduction to Routing



- Routing:
 - Discovery of paths between *first-hop router* and the *destination router*
- Path *routing metric*
 - Shortest path
 - Number of links
 - Sum of link delay
 - Lowest cost path
 - Sum of link costs
 - Highest throughput path
 - Maximum link capacity
 - Highest path reliability
 - Prob. of contact between path nodes in an ad-hoc mobile network



- Network model
 - Graph $G = (V, E)$
 - $V = \{u, v, w, x, y, z\}$
 - $E = \{(u, v), (u, x), (v, x), \dots\}$
 - Link cost $c(x, y)$
 - $c(x, y) = c(y, x)$
 - If $(x, y) \notin E$ then $c(x, y) = \infty$

Mountains & Minds

203


Introduction to Routing



What is the least cost path between u and z?

- Come up with a routing algorithm based on global knowledge of link costs
 - What happens when the network is very large?
 - What is the effect of frequent changes in network topology?

Link State (LS)



- Network model
 - Graph $G = (V, E)$
 - $V = \{u, v, w, x, y, z\}$
 - $E = \{(u, v), (u, x), (v, x), \dots\}$
 - Link cost $c(x, y)$
 - $c(x, y) = c(y, x)$
 - If $(x, y) \notin E$ then $c(x, y) = \infty$

Mountains & Minds

204

Dijkstra's Algorithm



```

1 Initialization:
2   N' = {u}
3   for all nodes v
4     if v adjacent to u
5       then D(v) = c(u,v)
6     else D(v) = ∞
7
8   Loop
9     find w not in N' such that D(w) is a minimum
10    add w to N'
11    update D(v) for all v adjacent to w and not in N' :
12      D(v) = min( D(v), D(w) + c(w,v) )
13    /* new cost to v is either old cost to v or known
14       shortest path cost to w plus cost from w to v */
15  until all nodes in N'

```

Update shortest path if
transitive cost through
added node is lower

Mountains & Minds


205

Dijkstra's algorithm example



Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					

So Dijkstra gives you the **cost** of the shortest path to each node. How would you find the **actual path** of nodes?



Mountains & Minds

206

Introduction to Routing




What is the least cost path between u and z?

- Come up with a routing algorithm based on global knowledge of link costs
 - What happens when the network topology changes?
 - What is the effect of frequent cost changes in network topology?
- Come up with a routing algorithm based on local knowledge of link costs?
 - Node x knows $c(x, v)$ for all its one hop neighbors
 - Node x can advertise its distance $d_x(y)$
 - Bellman-Ford equation:

$$d_x(y) = \min_{v \in V} [c(x, v) + d_v(y)]$$
 - What is the effect of frequent cost changes in network topology?

Link State (LS)

Distance Vector (DV)



- Network model
 - Graph $G = (V, E)$
 - $V = \{u, v, w, x, y, z\}$
 - $E = \{(u, v), (u, x), (v, x), \dots\}$
 - Link cost $c(x, y)$
 - $c(x, y) = c(y, x)$
 - If $(x, y) \notin E$ then $c(x, y) = \infty$

Mountains & Minds


207

Distance Vector (DV) Algorithm

```

1 Initialization:
2   for each destination y in N:
3      $D_x(y) = \infty$ 
4   for all neighbors y in N:
5      $D_x(y) = c(x, y)$ 
6   for all destination y in N:
7     send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to neighbors
  
```

Forwarding table num. rows equal to num. neighbors



```

8 while True:
9   wait for new distance vector from v
10  for each y in N:
11     $D_x(y) = \min_v \{c(x, v) + D_v(y)\}$ 
12    if  $D_x(y)$  change for any destination y:
13      send  $D_x = [D_x(y): y \text{ in } N]$  to all neighbors
  
```

Update shortest path if transitive cost through announcing node is lower

Mountains & Minds

208


DV Example

Node	Cost to	Node	Cost to	Node	Cost to
x	x y z	x	x y z	x	x y z
From		From		From	
x	0 2 7	x	0 2 3	x	0 2 3
y	∞ ∞ ∞	y	2 0 1	y	2 0 1
z	∞ ∞ ∞	z	3 1 0	z	3 1 0

Node	Cost to	Node	Cost to	Node	Cost to
y	x y z	y	x y z	y	x y z
From		From		From	
x	∞ ∞ ∞	x	0 2 7	x	0 2 3
y	2 0 1	y	2 0 1	y	2 0 1
z	∞ ∞ ∞	z	7 1 0	z	3 1 0

Node	Cost to	Node	Cost to	Node	Cost to
z	x y z	z	x y z	z	x y z
From		From		From	
x	∞ ∞ ∞	x	0 2 7	x	0 2 3
y	∞ ∞ ∞	y	2 0 1	y	2 0 1
z	7 1 0	z	3 1 0	z	3 1 0

$$\begin{aligned}
 D_x(z) &= \min_v \{c(x,v) + D_v(y)\} \\
 &= \min\{c(x,x) + D_x(z), \\
 &\quad c(x,y) + D_y(z)\} \\
 &= \min\{0+7, 2+1\} = 3
 \end{aligned}$$




209

DV Examples




Link cost decrease



- At time instant t_1 node y notices cost change and updates its D vector
- At t_2 z receives y's D vector and updates its $D_z(x)$ to 2
- At t_3 y receives z's D vector, makes no change to own D and the algorithm come to a *quiescent* state

Link cost increase




- At t_1 y notices cost change and

DV Poisoned Reverse



Link cost increase



- At t_1 y notices cost change and updates its $D_y(x)$ to 6
- At t_2 z receives y's D vector and updates its $D_z(x)$ to 7
 - Routing loop!
- At t_3 y receives z's D vector and updates its $D_y(x)$ to 8
 - Count to infinity problem!
- If z routes through y to x, z advertises *poisoned* $D_z(x) = \infty$ to y
- At t_1 y notices cost change and updates its $D_y(x)$ to 60
- At t_2 z receives y's D vector and **keeps** its $D_z(x)$ at 50
- At t_3 y receives z's D vector (no longer *poisoned*) and updates its $D_y(x)$ to 51
- Count to infinity problem not solved for routing loops with 3+ nodes...

Mountains & Minds

211

Link State vs. Distance Vector



- Mechanism
 - LS: centralized, synchronous
 - DV: distributed, asynchronous
- Message complexity
 - LS: $O(VE)$ messages sent
 - DV: exchange between neighbors only
- Speed of convergence
 - LS: $O(n \log n)$ algorithm
 - DV: convergence time varies
 - May get routing loops
 - Count-to-infinity problem
- Robustness: what happens if router malfunctions?
 - LS:
 - Node can advertise incorrect link cost
 - Each node computes only its own table
 - DV:
 - DV node can advertise incorrect path cost
 - Each node's table used by others
 - Errors propagate through the network

Mountains & Minds

212