



## Software Defined Networks (SDNs)

Mountains & Minds

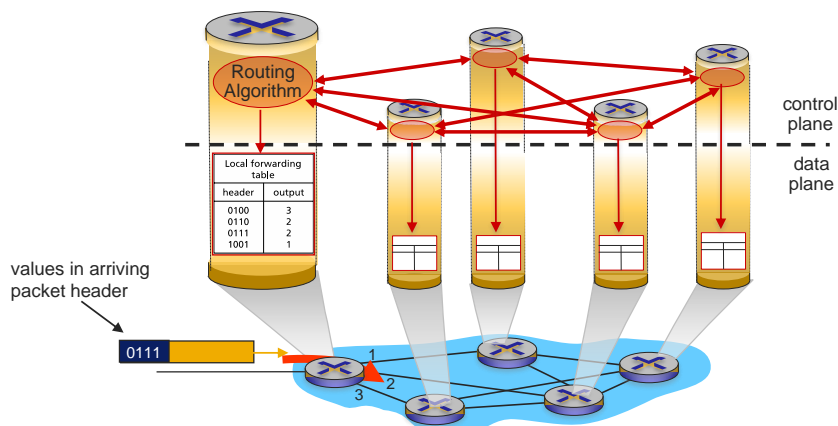
235

235

### Per-router control plane



Individual routing algorithm components *in each and every router* interact in the control plane



Mountains & Minds

236

236


# Traffic engineering



## Difficulties in traditional routing

**Q:** what if network operator wants u-to-z traffic to flow along uvwz?

**A:** need to define link weights so traffic routing algorithm computes routes accordingly



*Link weights are only control “knobs” in distributed routing protocols: wrong!*

Mountains & Minds

237

237


# Traffic engineering



## Difficulties in traditional routing

**Q:** what if network operator wants to split u-to-z traffic along uvwz and uxyz (load balancing)?

**A:** can't do it with – DV or LS mechanisms do not keep per flow state



Mountains & Minds

Network Layer: Control Plane

238


# Traffic engineering



## Difficulties in traditional routing

**Q:** what if  $w$  wants to route blue and red traffic differently?

**A:** can't do it (with destination based forwarding, and LS, DV routing)



Mountains & Minds


Network Layer: Control Plane

239

## Logically centralized control plane



A logically centralized controller interacts with local control agents (CAs) and installs forwarding rules (forwarding tables)



Mountains & Minds

240

240

# OpenFlow example




match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)

Example: datagrams from hosts h5 and h6 should be sent to h3 or h4, via s1 and from there to s2

match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)



Mountains & Minds

241

241

# OpenFlow example




match	action
IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(3)

What fields should SDN rules match?

What actions should routers be able take?

match	action
ingress port = 1 IP Src = 10.3.*.* IP Dst = 10.2.*.*	forward(4)

match	action
ingress port = 2 IP Dst = 10.2.0.3	forward(3)
ingress port = 2 IP Dst = 10.2.0.4	forward(4)

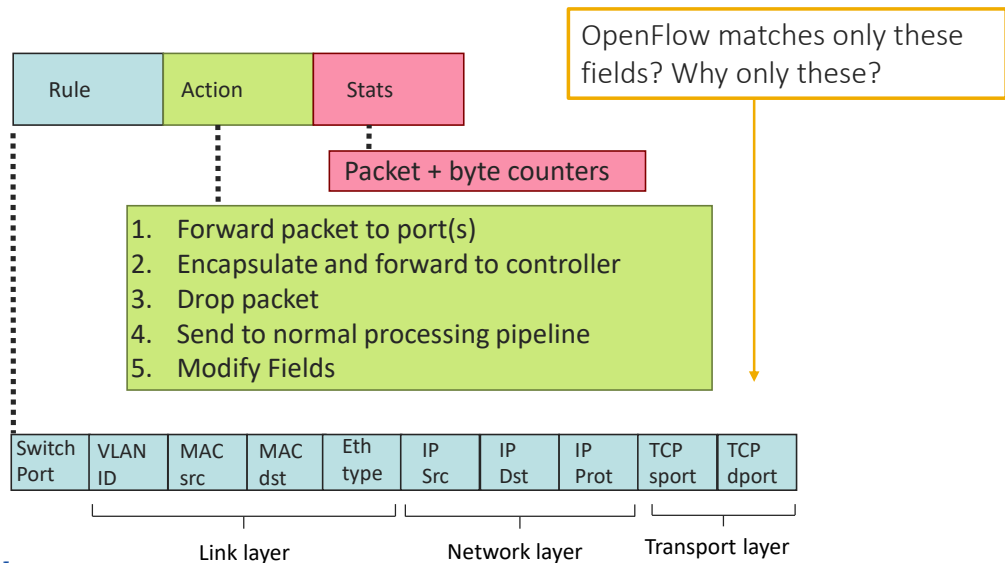


Mountains & Minds

242

242

# OpenFlow: Flow Table Entries



Mountains &amp; Minds


243

243


## Examples



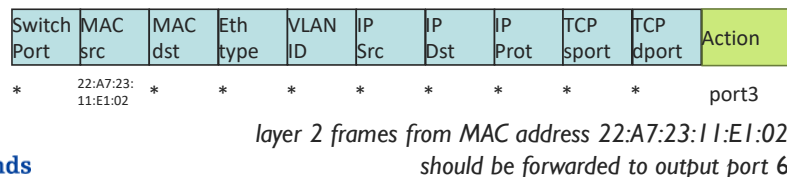
Destination-based forwarding:



Firewall:



Destination-based layer 2 (switch) forwarding:



Mountains &amp; Minds

244

244

# OpenFlow abstraction



match+action: unifies different kinds of devices

- Router
  - match: longest destination IP prefix
  - action: forward out a link
- Switch
  - match: destination MAC address
  - action: forward or flood
- Firewall
  - match: IP addresses and TCP/UDP port numbers
  - action: permit or deny
- NAT
  - match: IP address and port
  - action: rewrite address and port

Mountains & Minds

245

245


## SDN Implementation



- So we have switches and a logically centralized controller.

1. What is the functionality of the controller?

2. How would you organize that functionality architecturally into layers and functional units?



Mountains & Minds


248

248

## Analogy: mainframe to PC evolution



- Vertically integrated
- Closed, proprietary
- Slow innovation
- Small industry



- Horizontal
- Open interfaces
- Rapid innovation
- Huge industry



Mountains & Minds

249

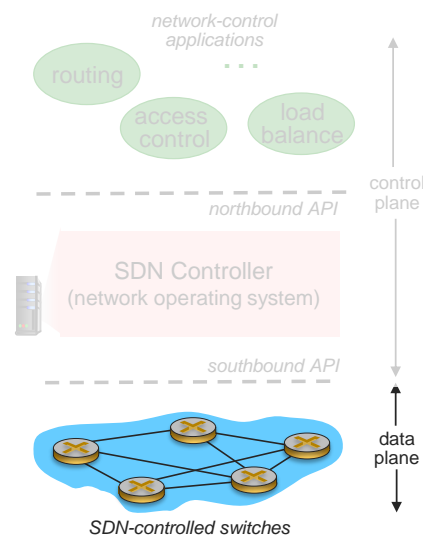
249

## SDN perspective: data plane switches



### Data plane switches

- Fast, simple, commodity switches implementing generalized data-plane forwarding in hardware
- Switch flow table computed, installed by controller
- API for table-based switch control (e.g., OpenFlow)
  - defines what is controllable and what is not
- Protocol for communicating with controller (e.g., OpenFlow)



Mountains & Minds

250

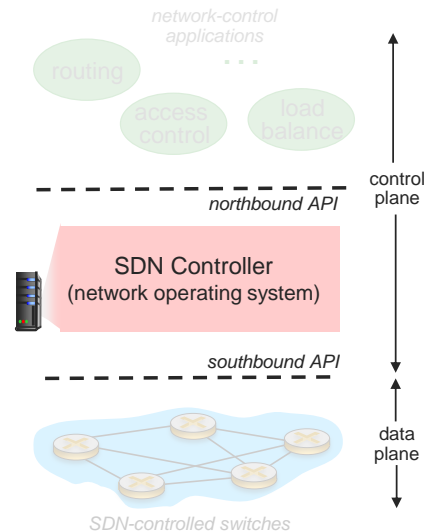
250

## SDN perspective: data plane switches



### SDN controller (network OS):

- Maintain network state information
- Interacts with network control applications “above” via northbound API
- Interacts with network switches “below” via southbound API
- Implemented as distributed system for performance, scalability, fault-tolerance, robustness



Mountains & Minds

251

251

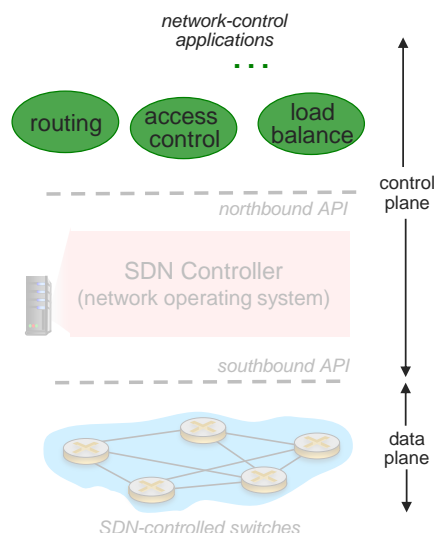
## SDN perspective: data plane switches



### Network-control apps:

- “Brains” of control: implement control functions using lower-level services, API provided by SDN controller
- Unbundled: can be provided by 3rd party: distinct from routing vendor, or SDN controller

Can we do anything? Transport (flow and congestion control, reliable delivery)?  
Application (caching)?



Mountains & Minds

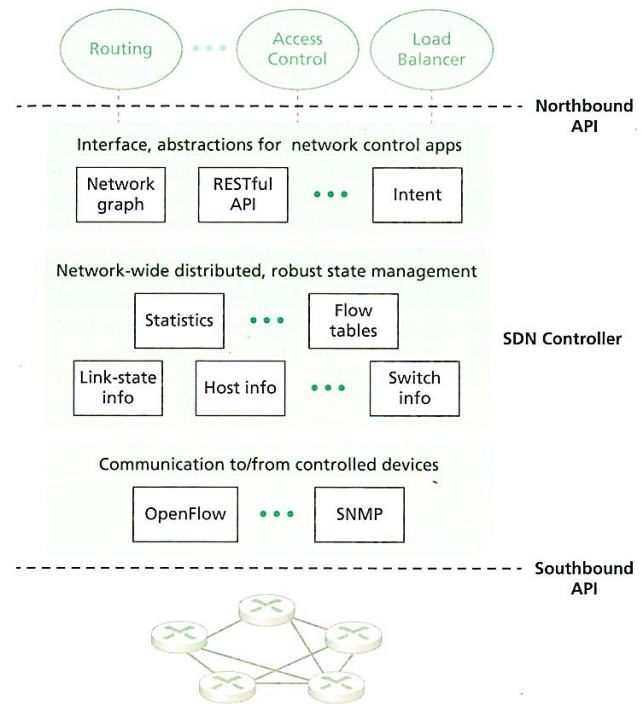
252

252



# SDN Organization

- Interface layer to network control apps:
  - Abstractions API
- Network-wide state management layer:
  - State of networks links, switches, services: a distributed database
- Communication layer:
  - Communicate between SDN controller and controlled switches



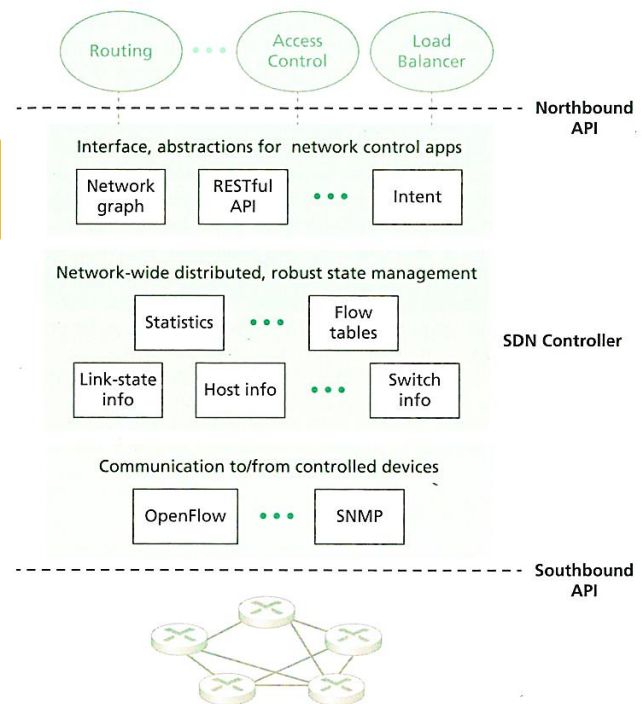
Mountains & Minds

253

# SDN Organization

What information needs is exchanged across the Southbound API?

- Southbound API
  - Switch to controller
    - Port status – includes hosts
    - Packet in – includes packets
    - Flow removed – timeout
  - Controller to switch
    - Read state – counters
    - Modify state – set forwarding table
    - Send packet – includes packet



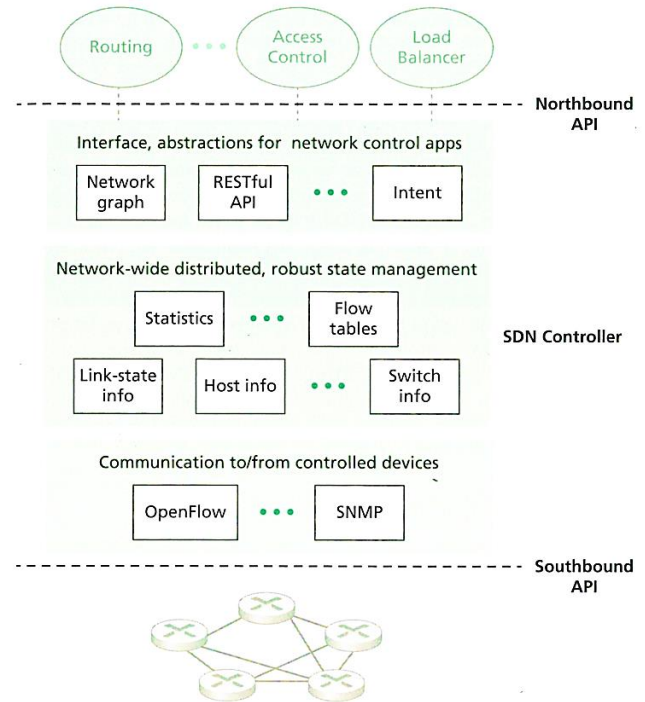
Mountains & Minds

254

# SDN Organization

What information needs to be exchanged across the Northbound API?

- Northbound API
  - Accept intents
  - Set flow tables explicitly
  - Query network graph
  - Subscribe to notifications

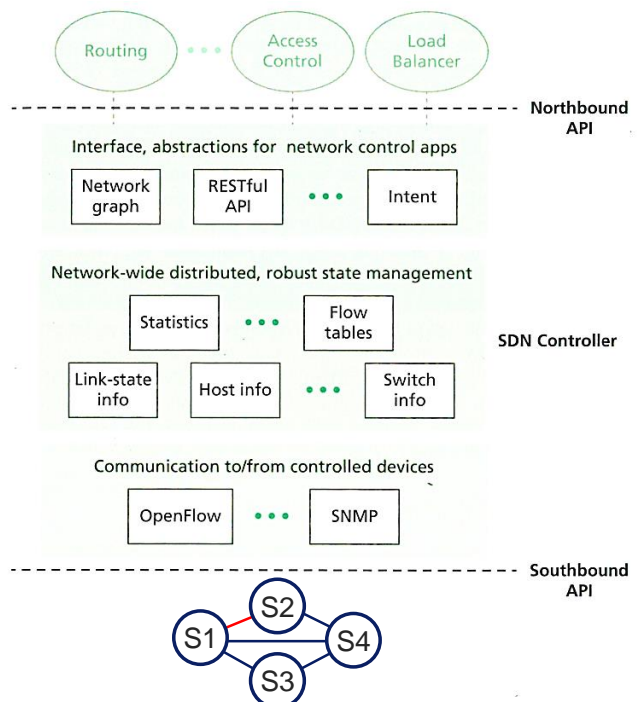


Mountains & Minds

255

## SDN: control/data plane interaction

What would be the set of steps within the SDN network when the link S1-S2 fails?

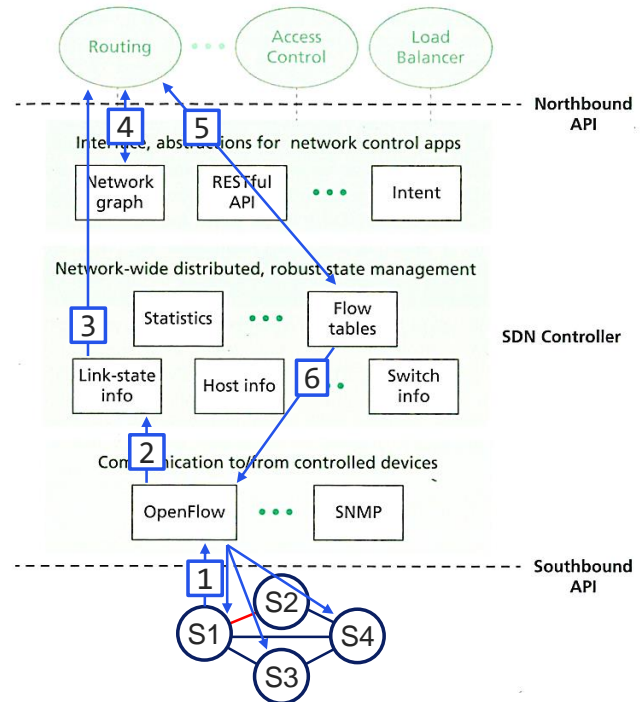


Mountains & Minds

256

# SDN: control/data plane interaction

1. S1, experiencing link failure using OpenFlow port status message to notify controller
2. SDN controller receives OpenFlow message, updates link status info
3. Dijkstra's routing algorithm application has previously registered to be called when ever link status changes. It is called.
4. Dijkstra's routing algorithm access network graph info, link state info in controller, computes new routes
5. Link state routing app interacts with flow-table-computation component in SDN controller, which computes new flow tables needed
6. Controller uses OpenFlow to install new tables in switches that need updating



Mountains & Minds


257

## Software Defined Networks (SDNs)



- SDN is a strategy for system, software, functionality development
  - Holistic L2-7 management platforms like OpenDaylight
  - Frenetic (Ox, Merlin, Nettle) declarative languages for network configuration
- Avoid dogma!
 

*"If you listen to the hype, software-defined networking and other software-defined technology will solve all data center problems and allow your computing gear to be an undifferentiated mass of identical nodes."*
- Limitations:
  - Inter-domain traffic?
  - Network virtualization and latency?
  - Computation speed vs. forwarding speed?
  - Table size explosion!



- Different controllers
  - OpenFlow – largest support for v.1
  - Floodlight – open source
  - OpenDaylight, OpenContrail – L2-7
  - Cisco APIC – proprietary extensions

Mountains & Minds

258

258